

操作系统研讨课

蒋德钧

Fall Term 2018-2019

email: jiangdejun@ict.ac.cn

office phone: 62601007



Lecture 2 A Simple Kernel

2018.09.26



Schedule

- Project 1 due
- Project 2 assignment



Project 1 Due

- P1 due
 - We test Tasks 2 and 3 for P1 due
 - Please compile your code, running the code on your board, and show the results to TA
 - Answer any questions we may ask



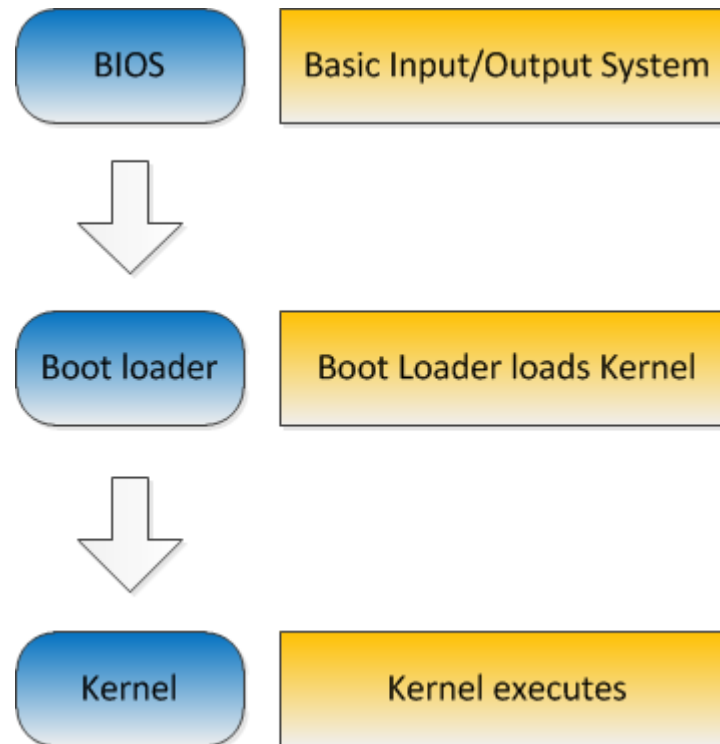
Project 1 Due

- P1 submission
 - Submit a compressed package named as “StudentNo.-YourName-P1”
 - Please includes
 - Source code
 - README to simply describe your code, e.g. which file is your work or how to run your code
 - Design document covering the questions in the design document template
 - Do not forget to submit before 23:55 TONIGHT



Project 2 – A Simple Kernel

- Booting procedure



Project 2 – A Simple Kernel

- Requirements
 - Write a simple kernel (non-preemptive)
 - Start a set of user processes and kernel threads
 - Perform context switches between processes and threads
 - Provide non-preemptive kernel support with context switch
 - Support basic mutex to allow BLOCK state of processes/threads



Project 2 – A Simple Kernel

- Requirements (Cont.)
 - Write a simple kernel (preemptive)
 - Provide preemptive kernel support with clock interrupt handler and priority based scheduling
 - Support system calls



Project 2 – A Simple Kernel

- A set of multiple tasks
 - Program codes under the *test* directory in start-code
 - Please refer to *test.c* for different groups of tasks
 - Fixed number of tasks for each test group
 - Allocate per-task state statically in main.c
 - STRONGLY suggest to first read the codes of different tasks to understand what they do



Project 2 – A Simple Kernel

- Process Control Block (PCB/TCB)
 - A data structure in OS kernel containing the information to manage a particular process/thread
 - Normally, kept in an area of memory protected from normal user access



Project 2 – A Simple Kernel

- Process Control Block (PCB/TCB)
 - Process status
 - Status of a process when it is suspended
 - Contents of registers, stack pointers etc.
 - Process scheduling info
 - e.g. priority



Project 2 – A Simple Kernel

- Start a task(process/thread)
 - Task type
 - User Process
 - Use faked SYSCALL to call kernel functions
 - But, in this project compiled with the kernel and share the same address with the kernel
 - Kernel thread
 - Task entry point
 - Function addresses
 - Please refer to *task_info* structure in start-code



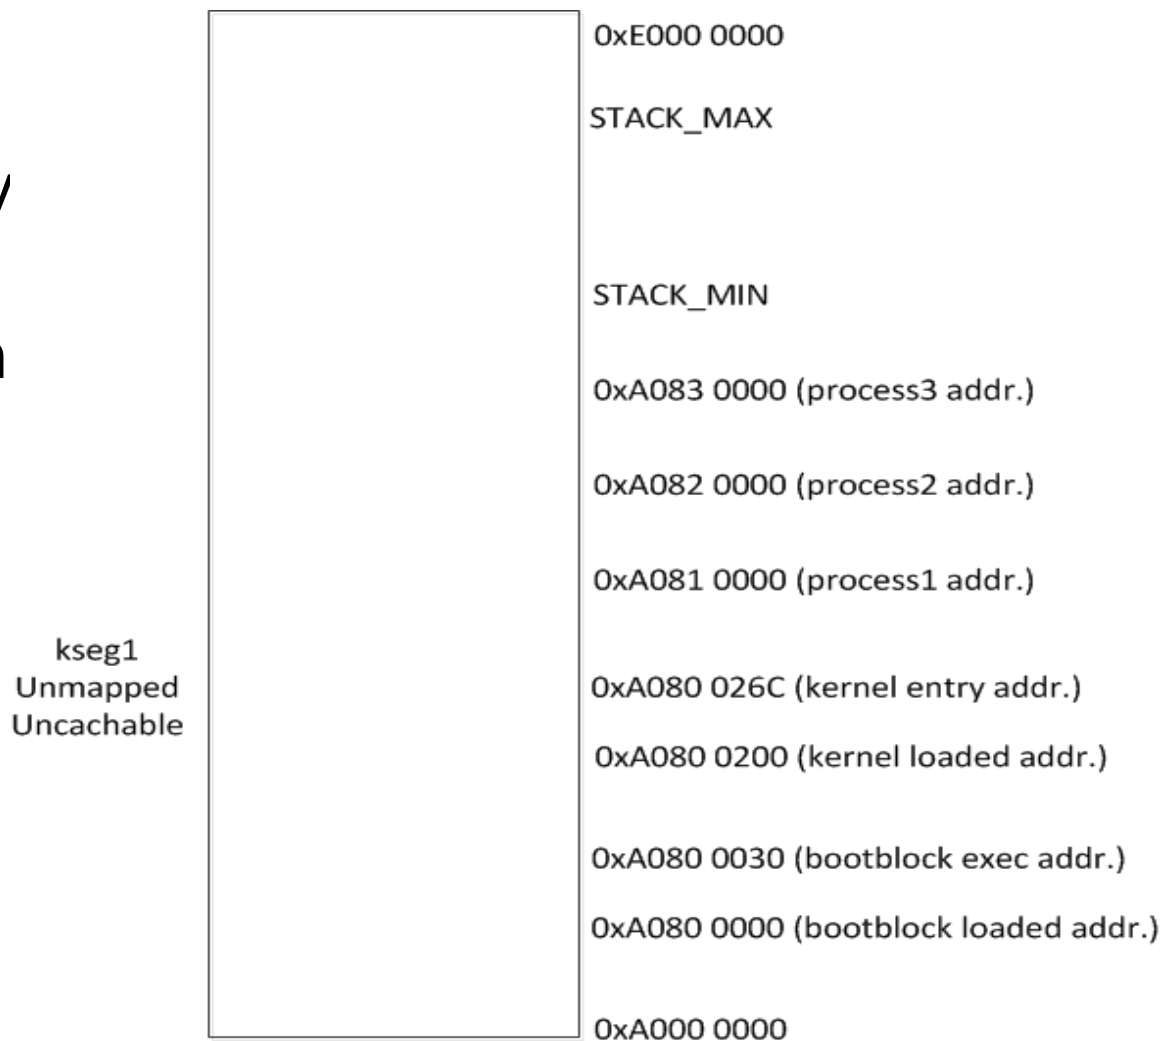
Project 2 – A Simple Kernel

- Start a task (process/thread)
 - Each task is associated with a PCB
 - Initialize PCB
 - Which registers should be set?
 - Where is the task located?
 - How to setup stack? Stack size?
 - Where is the PCB located?



Project 2 – A Simple Kernel

- Start a task
 - Possible memory layout
 - Decide your own STACK_MIN and STACK_MAX



Project 2 – A Simple Kernel

- Start a task
 - Scheduler: single task vs. multi-tasks
 - How to organize tasks being scheduled?
 - How to select the next task?
 - FIFO



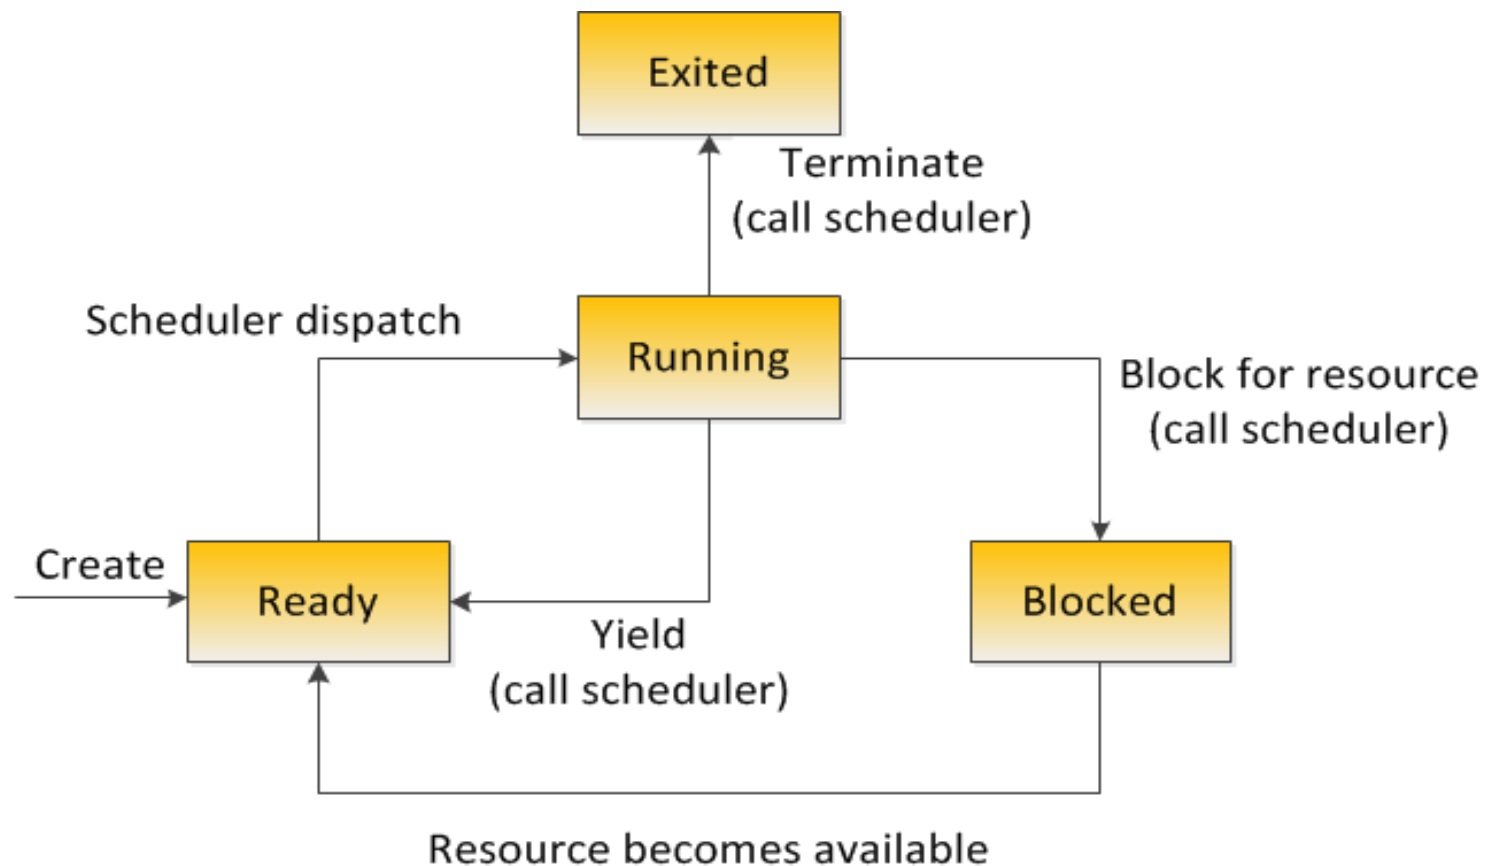
Project 2 – A Simple Kernel

- Start a task
 - Scheduler: first task vs. the following ones
 - do_scheduler()
 - Locate the PCB of the first task
 - Restore PCB
 - When the scheduler is invoked?
 - Non-preemptive kernel
 - Preemptive kernel



Project 2 – A Simple Kernel

- Scheduler (non-preemptive kernel)



Project 2 – A Simple Kernel

- Yield
 - An action to force a processor to release control of the current running thread
 - Place the current running thread to the end of the running queue
 - In this project, we call `do_scheduler()` to execute yield



Project 2 – A Simple Kernel

- Context switch
 - Save PCB
 - What kind of things to save
 - Registers → Memory
 - Restore PCB
 - Memory → Registers
 - `do_scheduler()`



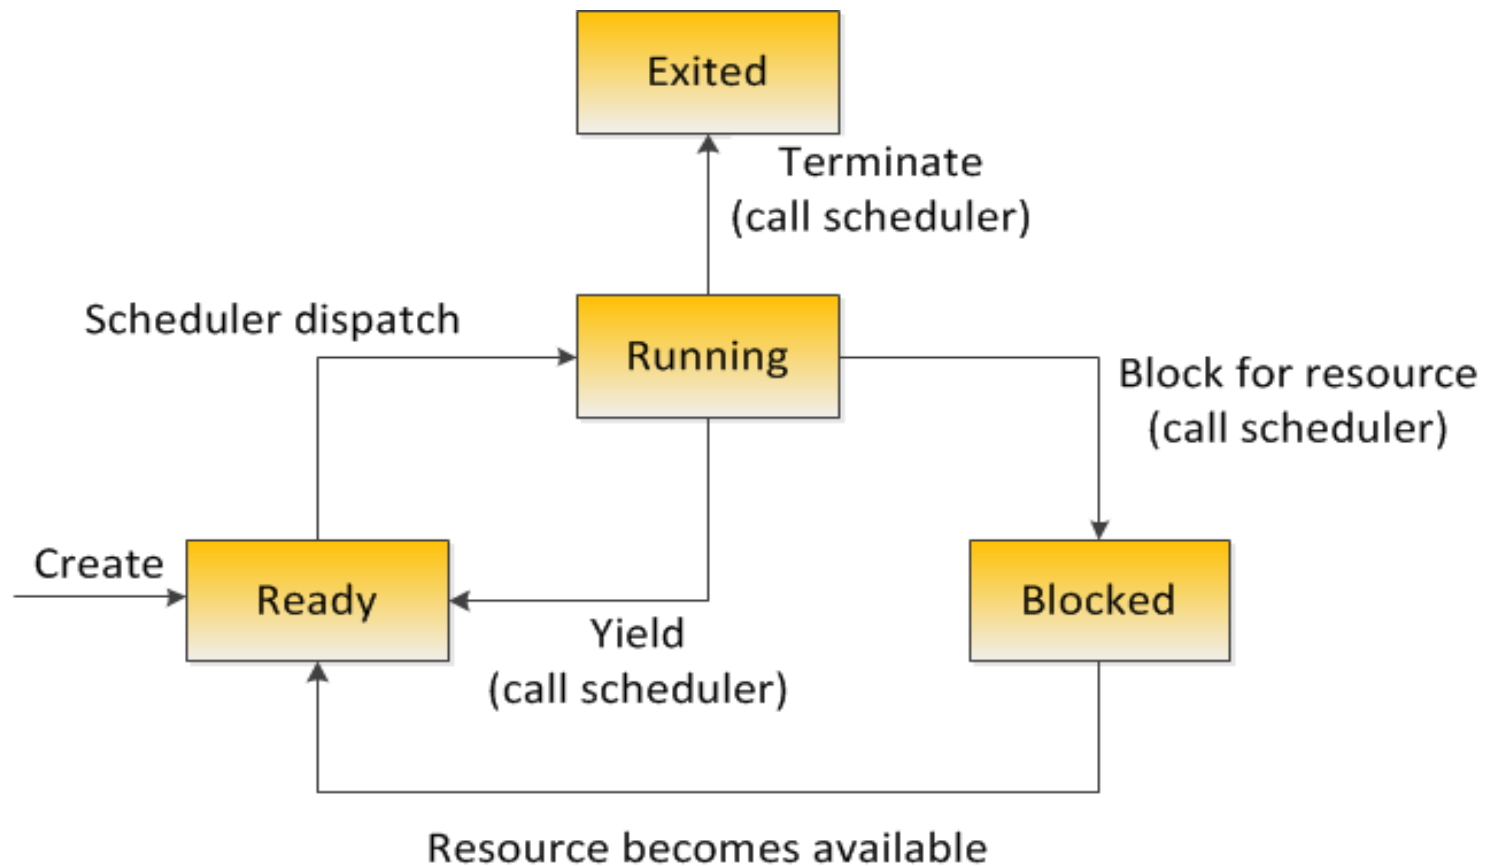
Project 2 – A Simple Kernel

- Create machine image
 - Combine binary of kernel, tasks together into a machine image using createimage in your P1



Project 2 – A Simple Kernel

- Mutex lock



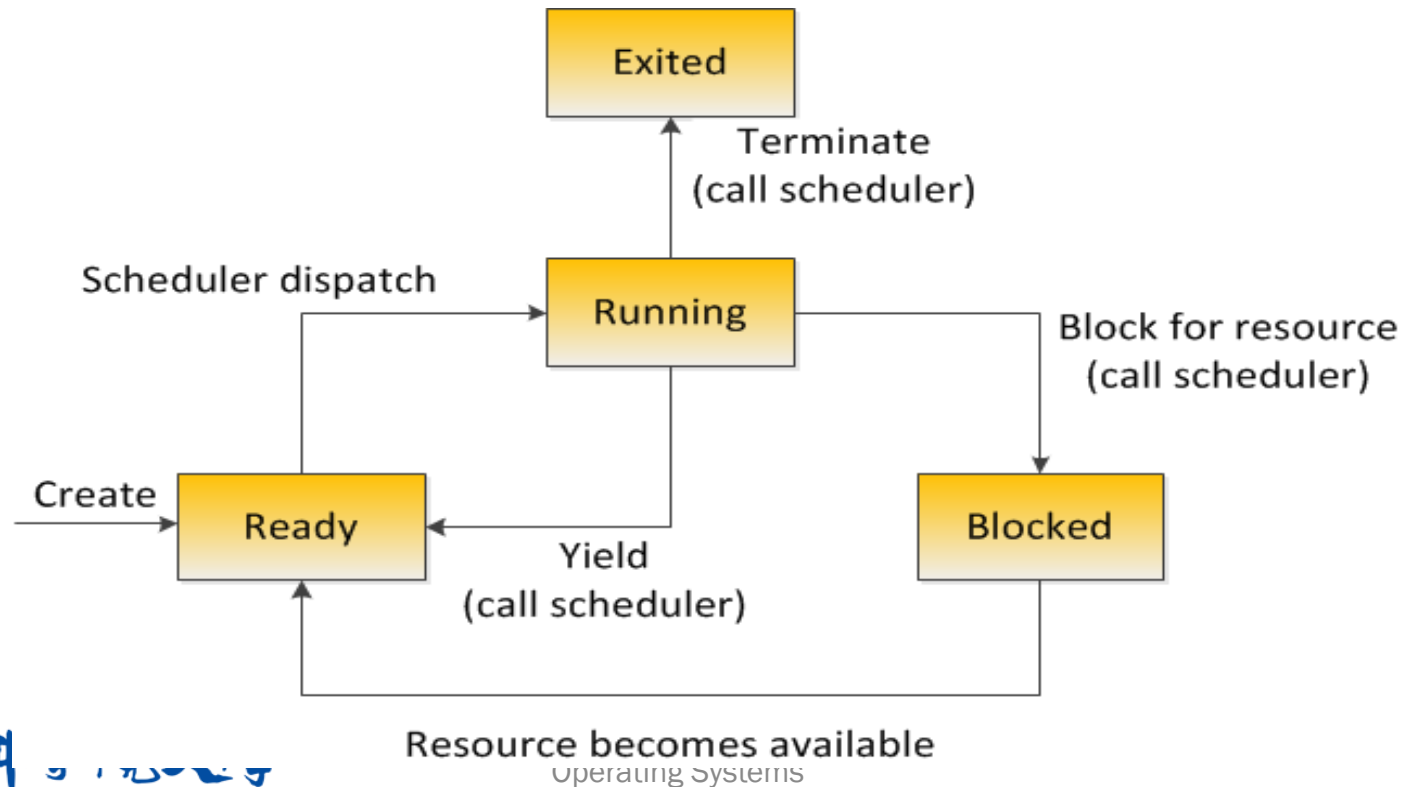
Project 2 – A Simple Kernel

- Mutex lock
 - What if no thread currently holds the lock?
 - What if the lock is currently held?
 - Implement lock-related functions
 - Manage tasks that do not acquire the lock
 - Same queue vs. different queues?



Project 2 A Simple Kernel

- Problems of Non-Preemptive kernel
 - How about the throughput of the operating system?



Project 2 A Simple Kernel

- Interrupt & Exception
 - A signal to the processor emitted by HW or SW indicating an event requiring immediate process
 - The processor responds by suspending its current running task, saving its state, and executing interrupt/exception handler
 - After the interrupt/exception handler finishes, the processor resumes normal activities




Project 2 A Simple Kernel

- Interrupt & Exception

Normal execution

Normal execution
with interrupt




Project 2 A Simple Kernel

- Hardware interrupt
 - A change in execution caused by an external event outside the processor
 - Clock for timesharing
 - I/O devices: disk, network, keyboard etc.
- Software interrupt
 - **System calls**: a user-programmed interrupt
- We do not handle other exceptions here
 - Segment fault, Overflow, Page fault etc.



Project 2 A Simple Kernel

- Handling interrupt
 - Save context
 - Determine what caused interrupt
 - Invoke specific routine based on type of interrupt
 - Restore context
 - Return from interrupt




Project 2 A Simple Kernel

- Handling interrupt
 - What if interrupt occurs while in interrupt handler?
 - ENTER_CRITICAL
 - Disable interrupt before actually handling the interrupt
 - Setting a label to indicate the interrupt is disabled
 - LEAVE_CRITICAL
 - Do not forget to re-enable interrupt after handling the interrupt



Project 2 A Simple Kernel

- Handling interrupt
 - ENTER_CRITICAL
 - Save context
 - Determine what caused interrupt
 - Invoke specific routine based on type of interrupt
 - LEAVE_CRITICAL
 - Restore context
 - Return from interrupt



Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - Coprocessor 0 (CP0) registers
 - Status register
 - IE bit: 1 enable interrupt; 0 disable
 - IM7 ~ IM0: control enable/disable different interrupts
 - IM7 bit: 1: enable Clock interrupt; 0: disable

31	28	27	26	25	24	23	22	21	20	19 16	15	8	7 5	4 3	2	1	0
CU (cu3:cu0)	0	FR	0	NO- FDIV	NO- FSQR	BEV	0	SR	0	IM7-IM0	0	KSU	ERL	EXL	IE		
4	1	1	1	1	1	1	1	1	4	8	3	2	1	1	1		




Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - Coprocessor 0 (CP0) registers
 - Status register
 - Refer to the initialization of interrupts in start code
 - Pay attention to the initialization of the register in YOUR PCB
 - We handle clock interrupt and syscall in this project



Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - Coprocessor 0 (CP0) registers
 - Cause register
 - EXCCODE: Exception type
 - Hardware interrupt / software interrupt
 - IP7 ~ IP0: indicating the interrupt type
 - IP7 bit: 1: Clock interrupt



Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - Operate CP0 registers
 - mfc0: move from Coprocessor 0
 - Loads data from a CP0 register into a CPU register
 - mtc0: move to Coprocessor 0
 - Stores data into a CP0 register
 - Use MIPS registers
 - \$k0, \$k1



Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - Coprocessor 0 (CP0): read-modify-write
 - Read the cp0 register into a CPU register
 - Modify the contents of the CPU register
 - Write the modified value back to the cp0 register
 - e.g. `mfc0 k0, CP0_STATUS`
`mtc0 k0, CP0_STATUS`



Project 2 A Simple Kernel

- Tips on implementing interrupt handler
 - What do you do in the clock interrupt handler?
 - How to deal with normal tasks?
 - How to deal with sleeping tasks?



Project 2 A Simple Kernel

- Process sleep()
 - Blocking sleep
 - Block the task when it calls sleep()
 - Use a separate queue to keep sleeping tasks
 - Wake up the task
 - When the timing reaches sleeping threshold of the task
 - About timing
 - A global counter recording the number of ticks
 - The counter increases in each clock interrupt
 - We provide wall time calculation functions in time.c



Project 2 A Simple Kernel

- Scheduler
 - Round robin
 - Priority based
 - Fairness
 - Think about what kind of information should be included in PCB if you want to do the above scheduler



Project 2 A Simple Kernel

- Step by step
 - Task 0: PLEASE read the guiding book and start code CAREFULLY
 - Task 1: start a set of kernel threads and support context switch as a non-preemptive kernel
 - Task 2: implement mutex lock to support BLOCK state



Project 2 A Simple Kernel

- Step by step
 - Task 3: implement a clock interrupt handler and priority based scheduler
 - Task 4: support syscalls and implement `sys_sleep`



Project 2 A Simple Kernel

- Requirement for design review (40 points)
 - What is PCB? What are included in PCB?
 - What need to be done for initializing tasks?
 - When is context switching in this project?
What need to be done for context switching?
 - How do you manage blocked tasks?



Project 2 A Simple Kernel

- Requirement for design review (40 points, cont.)
 - What is the workflow for handling interrupt (for both clock interrupt and syscalls)?
 - How do you implement the priority based scheduler?
 - When do you wake up the sleeping task?



Project 2 – A Simple Kernel

- Requirement for developing (60 points)
 - Start tasks and set PCBs: 5
 - Execute context switch without errors: 20
 - Implement the mutex lock: 5
 - Implement clock interrupt handler and syscall handler 20
 - Implement blocking sleep 5
 - Implement a priority-based scheduler 5



Project 2 – A Simple Kernel

- Bonus (1 point)
 - Support one thread to acquire multiple locks (at least two)
 - Support more than two threads to acquire a single lock
 - Design your own test cases and show the results



Project 2 – A Simple Kernel

- P2 schedule
 - P2 design review: 10th Oct.
 - No class on 17th Oct.
 - P2 due: 24th Oct.
- Final reminder
 - 重要的事情说三遍
 - Start early
 - Start early
 - Start early

