

操作系统研讨课

蒋德钧

Fall Term 2018-2019

email: jiangdejun@ict.ac.cn

office phone: 62601007



Lecture 5 Device Driver

2018.12.05



Schedule

- Project 4 due + Check P2、 P3
- Project 5 assignment



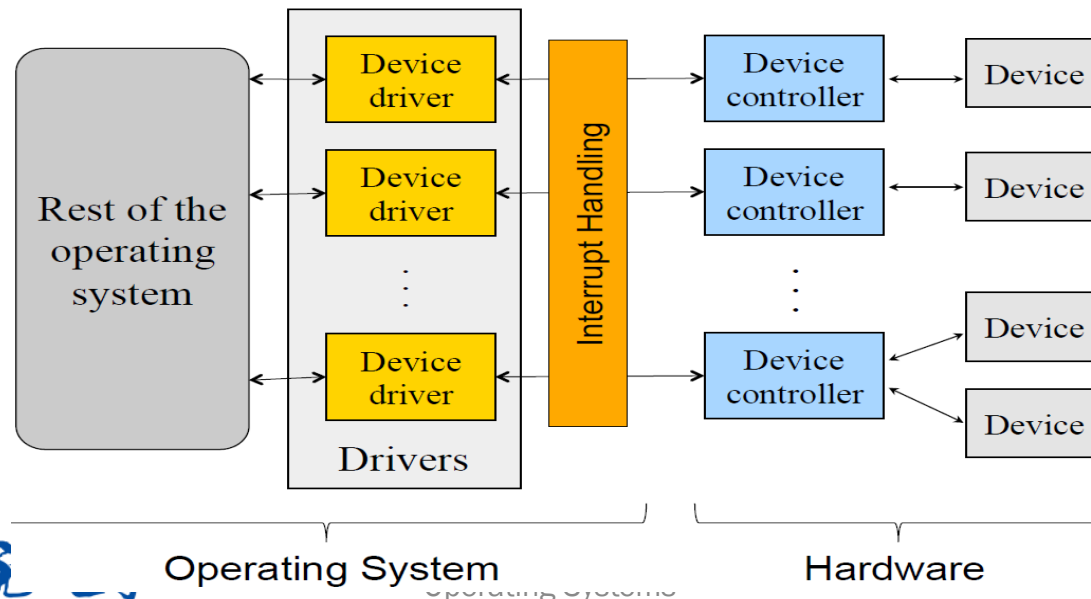
Project 5 Device Driver

- Requirement
 - Implement driver for network card
 - Setup MAC controller registers to allow network card initialization, and sending/receiving data without interrupt
 - Implement sending/receiving data with blocking
 - Implement network interrupt handler to serve receiving data



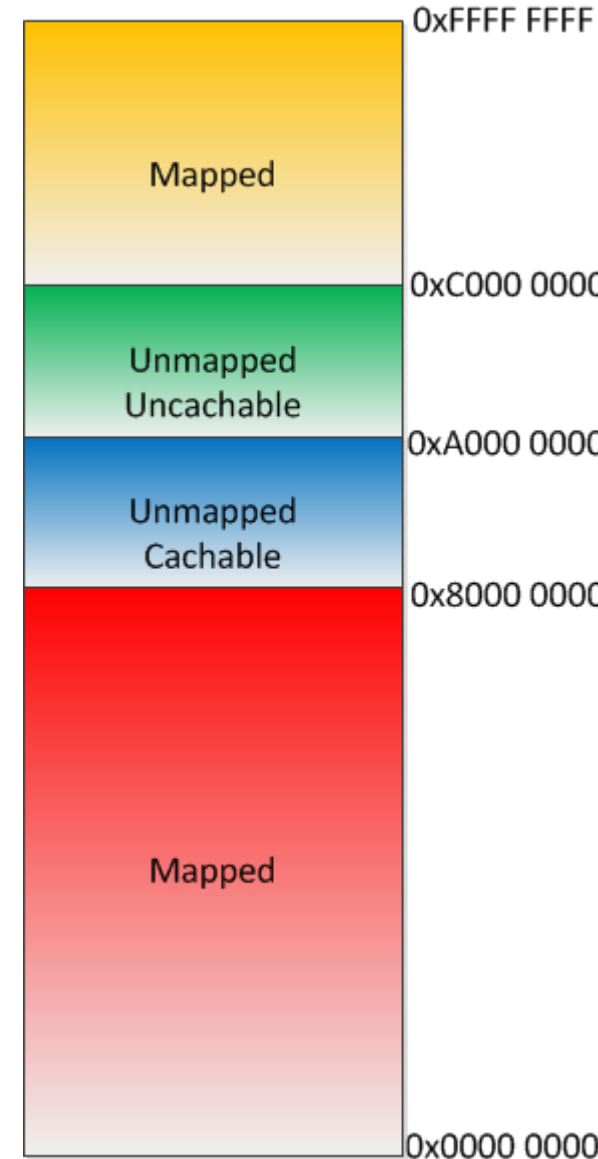
Project 5 Device Driver

- Device driver
 - Interact with device controller
 - Issue commands to device controller to finish data reads and writes, e.g. network and disk



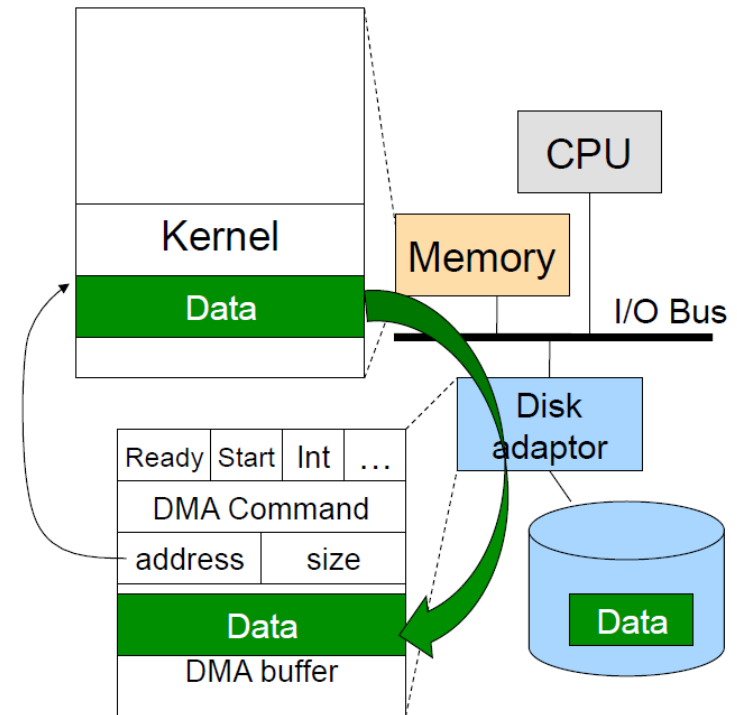
Project 5 Device Driver

- MAC controller
 - Include two parts of registers
 - MAC registers
 - Base address 0xbfe10000
 - Totally 462 registers
 - DMA registers
 - Base address 0xbfe11000
 - Totally 22 registers
 - Accessing a register
 - Base address + $i * 4$
 - i indicates i^{th} register (i starts from 0)



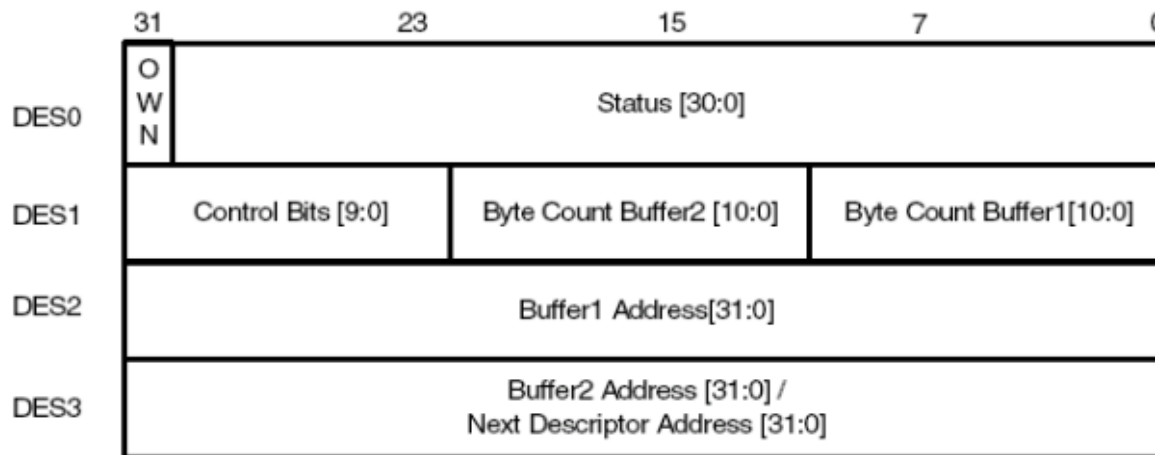
Project 5 Device Driver

- Direct Memory Access
 - Exchange data between host and device bypassing CPU
 - DMA transmit descriptor
 - DMA receive descriptor



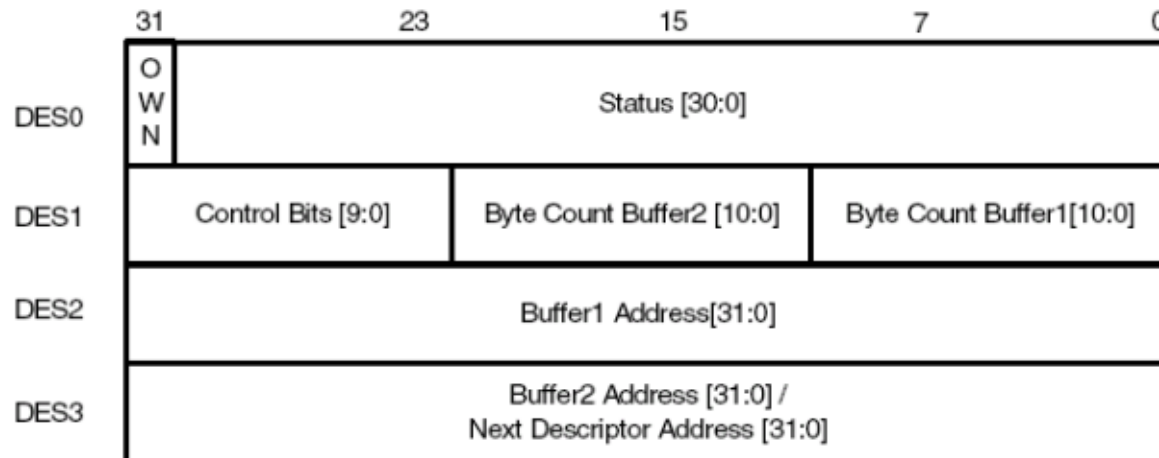
Project 5 Device Driver

- DMA descriptor
 - 16 bytes
 - Transmit descriptor for transmitting data
 - Receive descriptor for receiving data
 - One descriptor for one network packet
 - You are required to handle 64 packets
 - A circle linked list is expected to be used for multiple packets



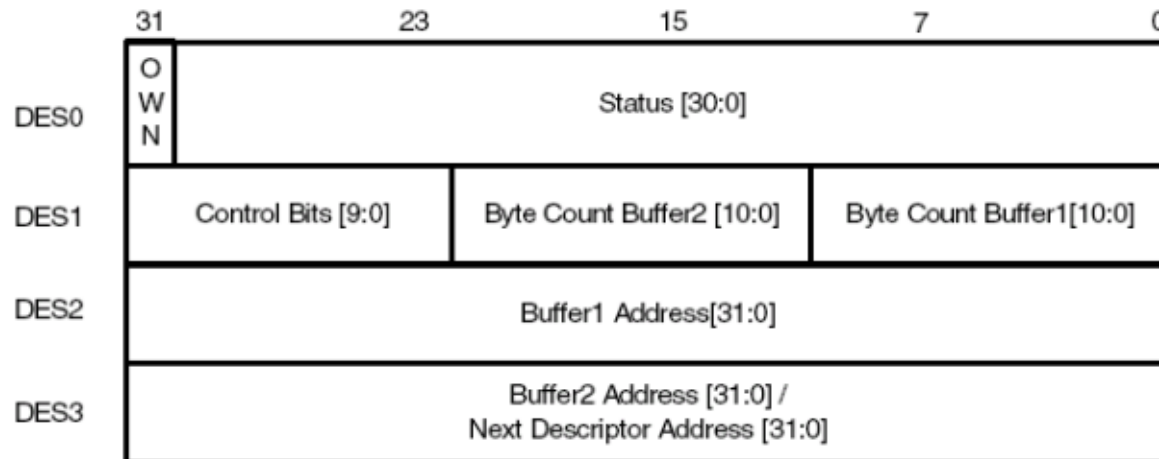
Project 5 Device Driver

- DMA descriptor
 - OWN flag
 - 0 indicates host has the rights to access DMA buffer
 - 1 indicates DMA controller has the rights to access DMA buffer
 - Remember to set this bit to 1 to allow DMA controller to Tx/Rx data



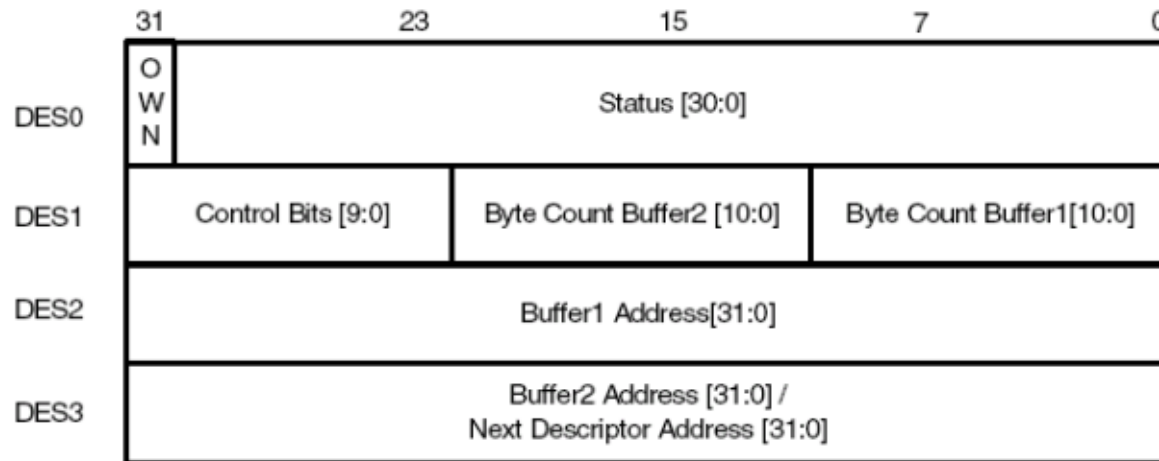
Project 5 Device Driver

- DMA descriptor
 - Status
 - Indicates different errors when transmitting or receiving data
 - Please refer to Fig. 4 and 7 for the meanings in the guiding book when you suffer from handling errors



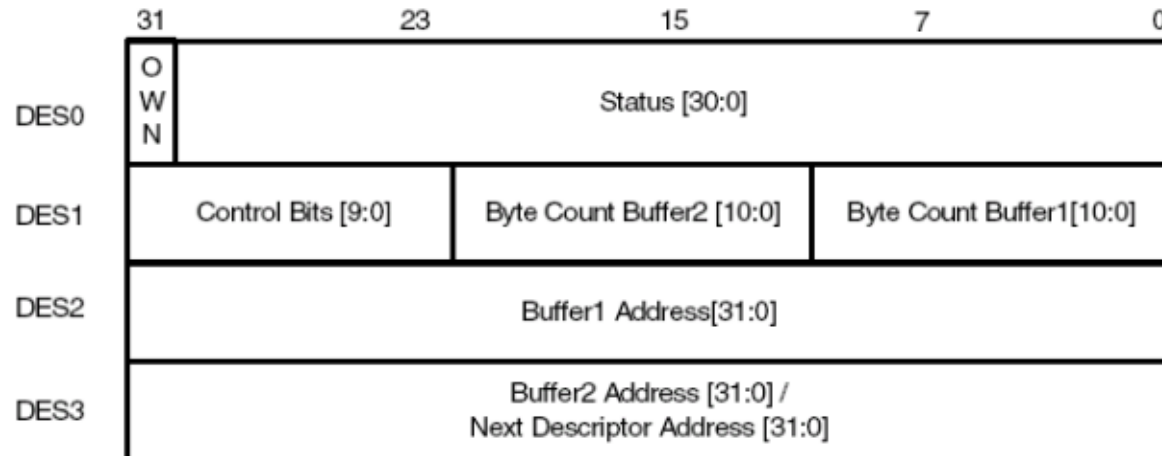
Project 5 Device Driver

- DMA descriptor
 - Control Bits
 - TER/RER (25) : Transmit/Receive End of Ring
 - 1 indicates this descriptor is the end of a descriptor ring
 - Please set this bit when initializing the descriptor ring



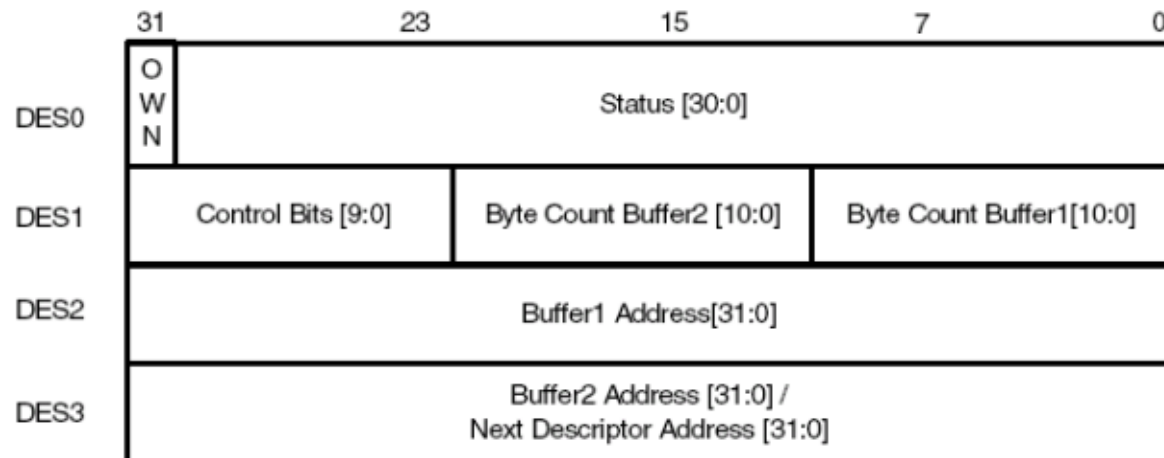
Project 5 Device Driver

- DMA descriptor
 - Control Bits
 - TCH/RCH (24) : Second Address Chained
 - 1 indicates the second buffer in this descriptor is linking the next descriptor
 - Please set this bit when initializing the descriptor ring



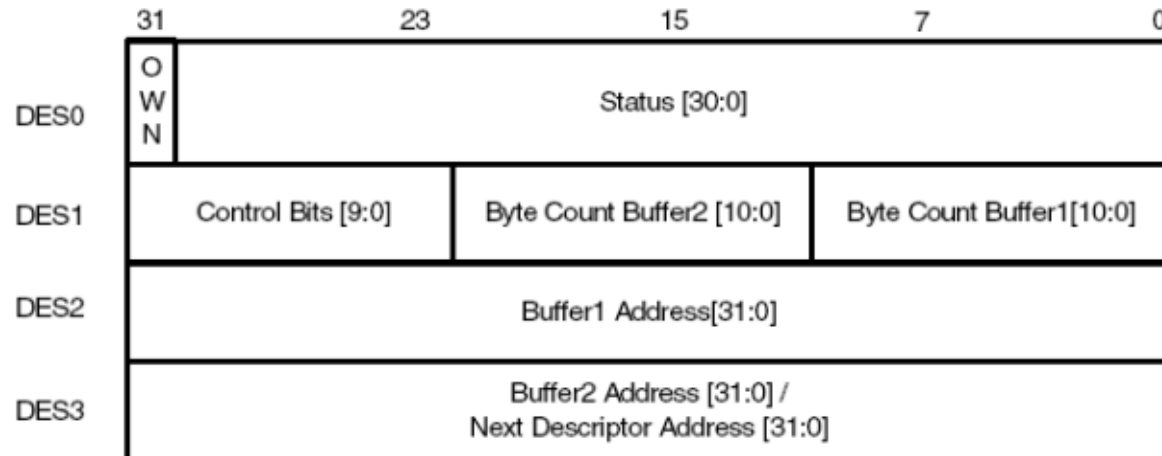
Project 5 Device Driver

- DMA descriptor
 - Control Bits
 - Byte Count Buffer1: the size of buffer1 in bytes, note that we have 1KB data to transmit/receive when filling buffer1
 - Byte Count Buffer2: set to 0 in this project



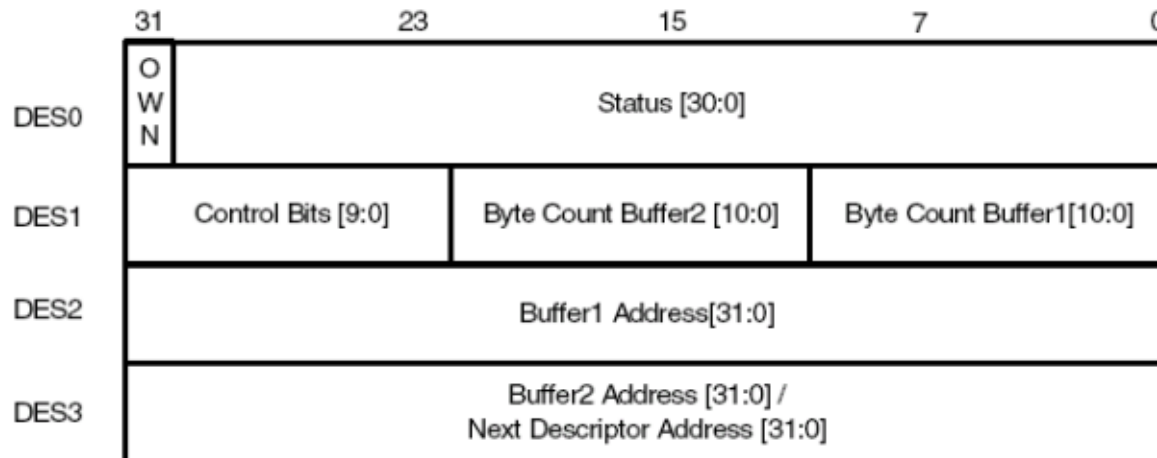
Project 5 Device Driver

- DMA descriptor
 - Buffer1 Address
 - The memory address of the starting address of data buffer
 - Note that, the memory address is in the unmapped area



Project 5 Device Driver

- DMA descriptor
 - Buffer2 Address
 - The memory address of next descriptor
 - Note that, we have 64 packets to transmit/receive respectively, and each packet is 1KB which needs a descriptor



Project 5 Device Driver

- Manipulating DMA descriptor
 - After filling DMA descriptor
 - You need to set the DMA register 4 and DMA register 3 to be the starting addresses of TX descriptor and RX descriptor, respectively
 - Before sending a packet
 - You need to write DMA register 1 with any value to trigger transmission
 - Before receiving a packet
 - You need to write DMA register 2 with any value to trigger receiving



Project 5 Device Driver


- Network packet
 - We already set the contents of each packet
 - The first 12 bytes refer to ethernet frame head
 - The following 20 bytes refer to IP head
 - Then the following 8 bytes refer to UDP head
 - The rest refer to data payload
 - Please do not change the packet contents UNLESS you are sure what you do

```
uint32_t buffer[PSIZE] = {0xffffffff, 0x5500ffff, 0xf77db57b,  
0x00450008, 0x0000d400, 0x11ff0040, 0xa8c073d8, 0x00e00101,  
0xe914fb00, 0x0004e914, 0x0000, 0x005e0001, 0x2300fb00,  
0x84b7f28b, 0x00450008, 0x0000d400, 0x11ff0040, 0xa8c073d8,  
0x00e00101, 0xe914fb00, 0x0801e914, 0x0000};
```



Project 5 Device Driver

- MAC interrupt handler
 - Cause register
 - exc-code: 0 refers to interrupt
 - IP3: 1 indicates device interrupt



Project 5 Device Driver

- MAC interrupt handler
 - 5 groups of registers correspond to device interrupts
 - Each group has 6 registers
 - `INTX_SR`
 - `INTX_EN`
 - `INTX_SET`
 - `INTX_CLR`
 - `INTX_POL`
 - `INTX_EDGE`



Project 5 Device Driver

- MAC interrupt handler
 - INTX_SR indicates the interrupt type (X is from 0 to 4)
 - e.g. INT0_SR refers to 32 types of interrupt
 - No. of MAC interrupt is 35
 - MAC interrupt is indicated by bit 3 in INT1_SR
 - MAC interrupt is enabled by bit 3 in INT1_EN
 - Please refer guiding book to understand how to set these bits using the register group



Project 5 Device Driver

- MAC interrupt handler
 - INT1_SR register
 - Address: 0xbfd01058
 - bit 3: 1 indicates MAC interrupt
 - INT1_EN register
 - Address: 0xbfd0105c
 - bit 3: 1 enables MAC interrupt



Project 5 Device Driver

- MAC interrupt handler
 - Check DMA RX descriptors to receive incoming data
 - Add this interrupt handler into *interrupt_helper*
 - Note that, you need to set corresponding IM bits when initializing interrupt



Project 5 Device Driver

- Step by step
 - Task 1: setup TX/RX DMA descriptors
 - Initialize DMA descriptors for transmitting and receiving data
 - Allocating data buffer for transmission and receiving, and fill given packets into transmit buffer
 - Set corresponding DMA registers to allow transmitting and receiving data
 - Implement syscalls for above functions
 - Note that, in this task, the receive thread continuously checks whether there is incoming data



Project 5 Device Driver

- Step by step
 - Task 2: implement MAC driver with blocking mode
 - Base on task1, but the receive thread blocks itself when there is no incoming data
 - The receive thread checks the incoming data in timer interrupt
 - The transmit thread does not change




Project 5 Device Driver

- Step by step
 - Task 3: handling MAC interrupt
 - Initialize and enable MAC interrupt
 - Implement MAC interrupt handler for receiving data
 - The transmit thread does not change



Project 5 Device Driver

- Task testing



Project 5 Device Driver

- Task testing
 - For transmitting data
 - Use wireshark (Windows) or tcpdump (Linux) to examine if the transmit thread successfully transmit data
 - Please read our reference documents
 - For receiving data
 - Use pktRxTx to transmit data from host to your development card to test receiving data



Project 5 Device Driver

- Requirements for design review (40 points)
 - What is DMA process? How do you set TX/RX descriptors? Which registers do you use to handle DMA process?
 - How many packets can you receive in task 1?
 - What is the procedure for receiving data in task 2?
 - What is the procedure of your MAC interrupt handler?
 - What is the procedure for transmitting data?



Project 5 Device Driver

- Requirements of developing (60 points)
 - Setup TX/RX descriptor and allow transmitting and receiving data (25)
 - Implement receiving data with blocking mode (15)
 - Implement MAC interrupt handler (20)



Project 5 Device Driver

- Bonus: receive large amount of data (2 points)
 - Use your MAC interrupt handler to receive large amount of data
 - pktTxRx can send data at 1MB/s ~ 2MB/s, let us send data for e.g. 10 seconds
 - Test the bandwidth of your receive thread, the expected bandwidth should be at least 1Mbit/s
 - Note that
 - your MAC interrupt handler should be robust enough to handle large amount of packets in a short time
 - You may need to handle packet loss



Project 5 Device Driver

- P5 schedule
 - P5 design review: 12th Dec.
 - Scheduled P5 due: 19th Dec.

